

Group 16

Joel Olausson & Sebastian Hermansson

We are running our tests on an 8th gen i7 laptop CPU and using 4 cores.

Assignment 1

Below you can see the benchmarking done on all jackknife functions that implement the different types of parallel maps. Only sJackknife and pmJackknife actually manage to outperform the sequential Jackknife function. Likely this is due to how the computer handles sparks depending on which parallelization techniques are used.

```
benchmarking jackknife
time                386.8 ms    (371.3 ms .. 400.3 ms)
                    1.000 R²    (0.999 R² .. 1.000 R²)
mean                421.8 ms    (407.5 ms .. 440.6 ms)
std dev             22.12 ms    (3.548 ms .. 29.46 ms)
variance introduced by outliers: 19% (moderately inflated)
```

```
benchmarking pJackknife
time                425.3 ms    (366.2 ms .. 457.9 ms)
                    0.998 R²    (0.994 R² .. 1.000 R²)
mean                437.2 ms    (428.8 ms .. 442.2 ms)
std dev             8.320 ms    (3.069 ms .. 11.47 ms)
variance introduced by outliers: 19% (moderately inflated)
```

```
benchmarking rJackknife
time                428.9 ms    (419.6 ms .. 446.6 ms)
                    1.000 R²    (1.000 R² .. 1.000 R²)
mean                427.7 ms    (424.9 ms .. 429.9 ms)
std dev             2.715 ms    (33.34 µs .. 3.238 ms)
variance introduced by outliers: 19% (moderately inflated)
```

```
benchmarking sJackknife
time                341.6 ms    (332.2 ms .. 349.0 ms)
                    1.000 R²    (1.000 R² .. 1.000 R²)
mean                344.2 ms    (342.7 ms .. 346.0 ms)
std dev             2.105 ms    (907.3 µs .. 2.813 ms)
variance introduced by outliers: 19% (moderately inflated)
```

```
benchmarking pmJackknife
time                365.9 ms    (75.56 ms .. 535.9 ms)
                    0.934 R²    (0.779 R² .. 1.000 R²)
mean                442.5 ms    (394.8 ms .. 469.7 ms)
std dev             45.29 ms    (5.182 ms .. 58.01 ms)
variance introduced by outliers: 23% (moderately inflated)
```

Below it is shown how jackknife run with our rmap implementation compares to the given parMap from the Strategies library. As is shown our rmap does not perform as well as parMap. If we compare parJackknife with the sJackknife implementation from the previous image it seems much more similar in performance but both of them perform significantly better than sJackknife.

```
benchmarking rJackknife
time                431.1 ms    (369.8 ms .. 466.8 ms)
                   0.998 R²    (0.996 R² .. 1.000 R²)
mean                456.1 ms    (437.2 ms .. 485.8 ms)
std dev             29.25 ms    (5.332 ms .. 37.39 ms)
variance introduced by outliers: 19% (moderately inflated)

benchmarking parJackknife
time                358.7 ms    (351.2 ms .. 366.2 ms)
                   1.000 R²    (1.000 R² .. 1.000 R²)
mean                354.2 ms    (351.1 ms .. 356.0 ms)
std dev             3.042 ms    (722.1 µs .. 4.096 ms)
variance introduced by outliers: 19% (moderately inflated)
```

Assignment 2

We implemented a mergesort and a sum function on our higher-order divide and conquer algorithm.

The mergesort had a little worse performance than the sequential merge sort before we set granularity. Through testing granularity by setting the base case list size to various numbers we determined that our best performance was achieved by setting it to 8. Below you can see the sequential as well as the parallel version. With optimal granularity, the divide-and-conquer implementation performs better than the sequential.

```
benchmarking mergesort
time                3.315 ms    (3.254 ms .. 3.362 ms)
                   0.998 R²    (0.997 R² .. 0.999 R²)
mean                3.451 ms    (3.399 ms .. 3.525 ms)
std dev             213.1 µs    (152.7 µs .. 323.3 µs)
variance introduced by outliers: 40% (moderately inflated)

benchmarking pmergesort
time                2.276 ms    (2.240 ms .. 2.331 ms)
                   0.995 R²    (0.991 R² .. 0.998 R²)
mean                2.259 ms    (2.221 ms .. 2.295 ms)
std dev             126.9 µs    (104.6 µs .. 155.8 µs)
variance introduced by outliers: 40% (moderately inflated)
```

The sum performed terribly without any granularity control.

```
benchmarking sum
time           39.55 µs    (38.97 µs .. 40.35 µs)
               0.998 R²    (0.996 R² .. 0.999 R²)
mean          40.89 µs    (40.27 µs .. 41.68 µs)
std dev       2.558 µs    (2.113 µs .. 3.040 µs)
variance introduced by outliers: 67% (severely inflated)

benchmarking psum
time           2.515 ms    (2.442 ms .. 2.601 ms)
               0.996 R²    (0.993 R² .. 0.999 R²)
mean          2.430 ms    (2.406 ms .. 2.466 ms)
std dev       94.86 µs    (73.79 µs .. 143.9 µs)
variance introduced by outliers: 23% (moderately inflated)
```

Here we can see it performing about 64 times worse than the regular sequential sum function. The overhead introduced by creating sparks and the required logic for divide and conquer is causing the computer to spend most of its time on things completely unrelated to actually adding numbers together.

```
benchmarking sum
time           46.80 µs    (46.35 µs .. 47.32 µs)
               0.998 R²    (0.997 R² .. 0.999 R²)
mean          48.67 µs    (47.95 µs .. 50.18 µs)
std dev       3.349 µs    (2.568 µs .. 4.351 µs)
variance introduced by outliers: 70% (severely inflated)

benchmarking psum
time           485.1 µs    (472.0 µs .. 498.4 µs)
               0.996 R²    (0.992 R² .. 0.999 R²)
mean          475.5 µs    (470.6 µs .. 481.2 µs)
std dev       18.63 µs    (13.06 µs .. 27.87 µs)
variance introduced by outliers: 33% (moderately inflated)
```

By introducing granularity where the normal sum function is run on lists that are smaller than 100 we can see a large improvement. Likely because the computer spends much less time on divide-and-conquer logic and spark creation and can instead use its time to actually add numbers.

Assignment 3

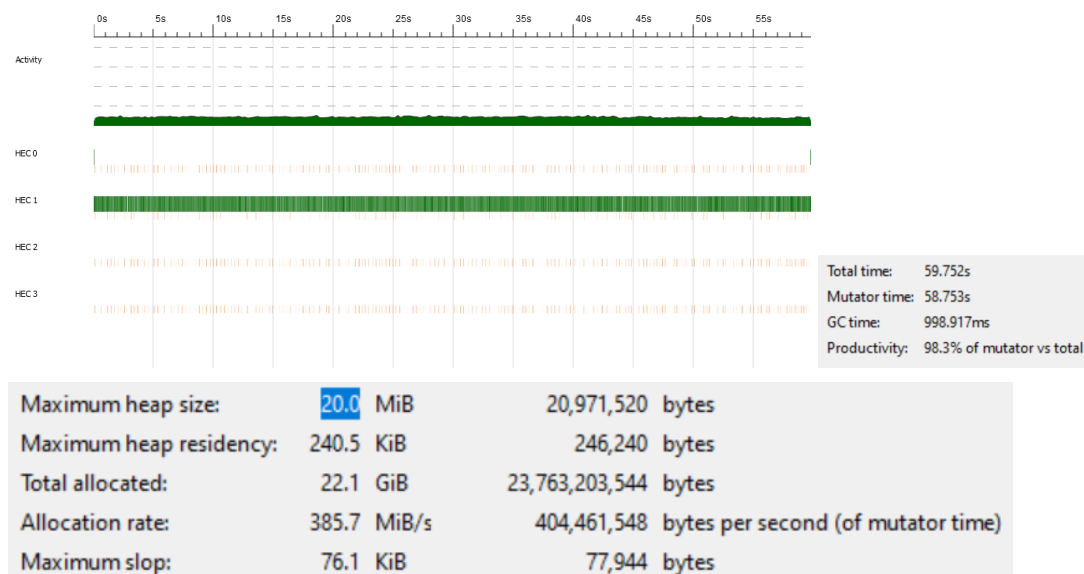
1

In short, `parBuffer` is a function that takes two arguments, an integer, let us call it “n”, and a strategy, that we will call “s”. The function causes evaluation on the inputs to be done with a limitation on available sparks to n at any one time. This makes the program less resource intensive on memory due to fewer sparks being on a waiting list and avoids too much parallelism which creates unnecessary overhead.


2

We chose to parallelize the sudoku solver. Essentially what we did was replace the `map` function with parallelized ones utilizing `parBuffer`, `parList` and `parListChunk` as well as a sequential one for comparison. We are able to see significant differences in how these map implementations handle the lists they are working with and how they subsequently handle sparks.

Sudoku running the regular `map` can be seen in threadscope only utilizing a single core. In total, it takes about a minute to run.




The sudoku using the `parBuffer` version of the `map` runs on all cores in parallel, however, the amount of work being done in parallel is rather low, as can be seen from the height of the activity reading. We can also see the spark pool which in this case shows that for the first 5 seconds HEC1 is the only one with sparks in its spark pool. After, that is instead the case for HEC2. This only happens for the version that restricts the buffer to 20 sparks. During the whole process, the number of sparks never exceeds 20. The execution time is about half that of the sequential `map`.



Maximum heap size:	21.0 MiB	22,020,096 bytes
Maximum heap residency:	437.1 KiB	447,560 bytes
Total allocated:	22.5 GiB	24,188,599,744 bytes
Allocation rate:	724.1 MiB/s	759,261,851 bytes per second (of mutator time)
Maximum slop:	90.8 KiB	92,936 bytes

With the buffer limited to 50, we see similar results in productivity and execution time. However, the number of sparks is consistently lower throughout the execution except for right in the beginning.




Maximum heap size:	21.0 MiB	22,020,096 bytes
Maximum heap residency:	508.8 KiB	520,984 bytes
Total allocated:	22.5 GiB	24,185,332,016 bytes
Allocation rate:	731.6 MiB/s	767,168,275 bytes per second (of mutator time)
Maximum slop:	95.1 KiB	97,352 bytes

For the 100-sized buffer there is once again a very similar execution time. The number of sparks once again rises quickly at the start but then ends up consistently even lower than that of the 50-sized buffer.



Maximum heap size:	21.0 MiB	22,020,096 bytes
Maximum heap residency:	620.1 KiB	634,976 bytes
Total allocated:	22.5 GiB	24,183,406,656 bytes
Allocation rate:	716.9 MiB/s	751,749,020 bytes per second (of mutator time)
Maximum slop:	106.4 KiB	108,976 bytes


The parList version of the map performs much worse. It starts creating many sparks and then resolves the sparks over time which is why we can see that the number of sparks goes down. This uses much more memory at a time than the parBuffer version which would be bad if we were limited. Online the buffered version this one overflows. It also stops running in parallel after around 7 seconds which is likely why it performs worse.



Maximum heap size:	117.0 MiB	122,683,392 bytes
Maximum heap residency:	32.3 MiB	33,863,440 bytes
Total allocated:	22.2 GiB	23,856,540,816 bytes
Allocation rate:	485.6 MiB/s	509,159,396 bytes per second (of mutator time)
Maximum slop:	1.4 MiB	1,444,504 bytes

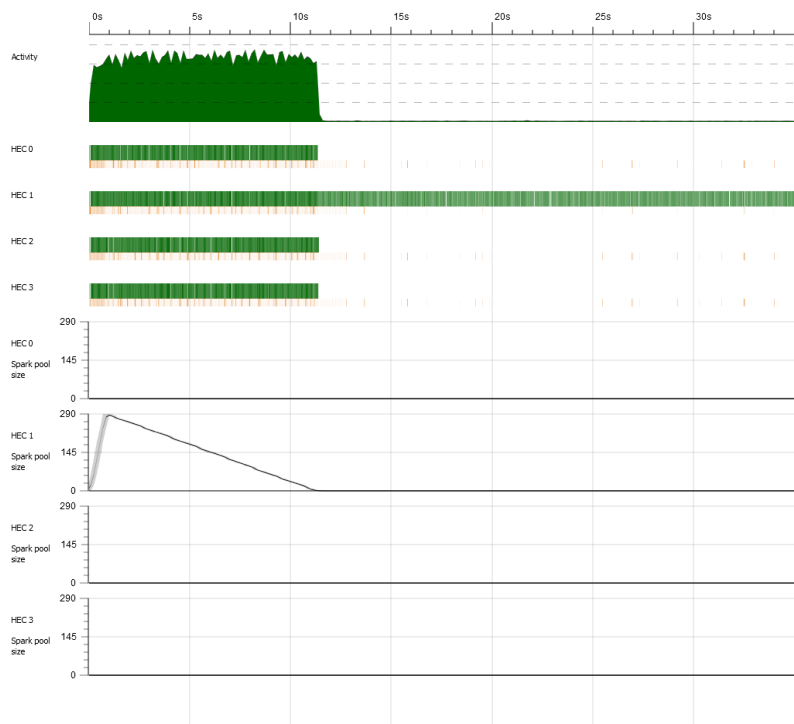
HEC	Total	Converted	Overflowed	Dud	GCed	Fizzled
Total	16000	7623	185	0	0	0
HEC 0	0	2204	0	0	0	0
HEC 1	16000	1033	185	0	0	0
HEC 2	0	2163	0	0	0	0
HEC 3	0	2223	0	0	0	0

For parListChunk the performance is very similar to that of parBuffer. For chunks set to 50, 100, and 200 neither the execution time, heap sizes nor any other factors really change. This could be because the list that is being worked on is too small for there to be a large difference between these list sizes. Similar to parList this version creates many sparks in the beginning and then resolves them but it continues running in parallel for 10 seconds which leads to better performance.




Total time: 34.870s
 Mutator time: 34.128s
 GC time: 742.729ms
 Productivity: 97.9% of mutator vs total

Maximum heap size:	115.0 MiB	120,586,240 bytes
Maximum heap residency:	31.5 MiB	33,064,712 bytes
Total allocated:	22.2 GiB	23,850,764,016 bytes
Allocation rate:	661.6 MiB/s	693,695,453 bytes per second (of mutator time)
Maximum slop:	2.3 MiB	2,403,104 bytes




Total time: 35.183s
 Mutator time: 34.382s
 GC time: 800.409ms
 Productivity: 97.7% of mutator vs total

Maximum heap size:	115.0 MiB	120,586,240 bytes
Maximum heap residency:	31.3 MiB	32,844,824 bytes
Total allocated:	22.2 GiB	23,850,495,888 bytes
Allocation rate:	656.4 MiB/s	688,257,471 bytes per second (of mutator time)
Maximum slop:	2.5 MiB	2,605,432 bytes



3

The chunked and buffered version of map runs similarly to that of parBuffer. It has a relatively continuous buffer size. Its heap size is much smaller than the parListChunk version, more similar to parBuffer. It has the advantage of avoiding too much parallelization that creates unnecessary overhead. A larger list chunk helps with performance in our testing. Activity is rather low throughout, however.



Total time:	32.642s
Mutator time:	31.851s
GC time:	791.423ms
Productivity:	97.6% of mutator vs total

Maximum heap size:	21.0 MiB	22,020,096 bytes
Maximum heap residency:	599.0 KiB	613,424 bytes
Total allocated:	22.5 GiB	24,189,055,064 bytes
Allocation rate:	724.3 MiB/s	759,452,382 bytes per second (of mutator time)
Maximum slop:	92.5 KiB	94,752 bytes